


More certified verdicts need not mean fewer verdict errors

20% budget · release order · Fixed 50/50 allocation minus random criterion sampling

Each bar is the mean of 31 paired seed differences at equal actual query counts. Additive counts; column scales differ.


(a) RuVerBench DR · Gemini 3.1 Pro

N = 284 tasks; 323 queried criteria




(b) JudgmentBench · GPT-5.4 · full


N = 1,539 annotations; 4,697 queried criteria



(c) Eager-update error decomposition



(d) Eager-update error decomposition



Negative residual change = fewer errors. Initial correctness, false passes and false fails are defined against the fixed released reference.

Components add before rounding. These are paired arithmetic means (seeds 1-31), not marginal medians or population effect estimates.