


Update effects across all budgets — DR Gemini 3.1 Pro

N = 284 analysis units; M = 1,615 criteria • shared scale for first two columns within this page; scales vary by page



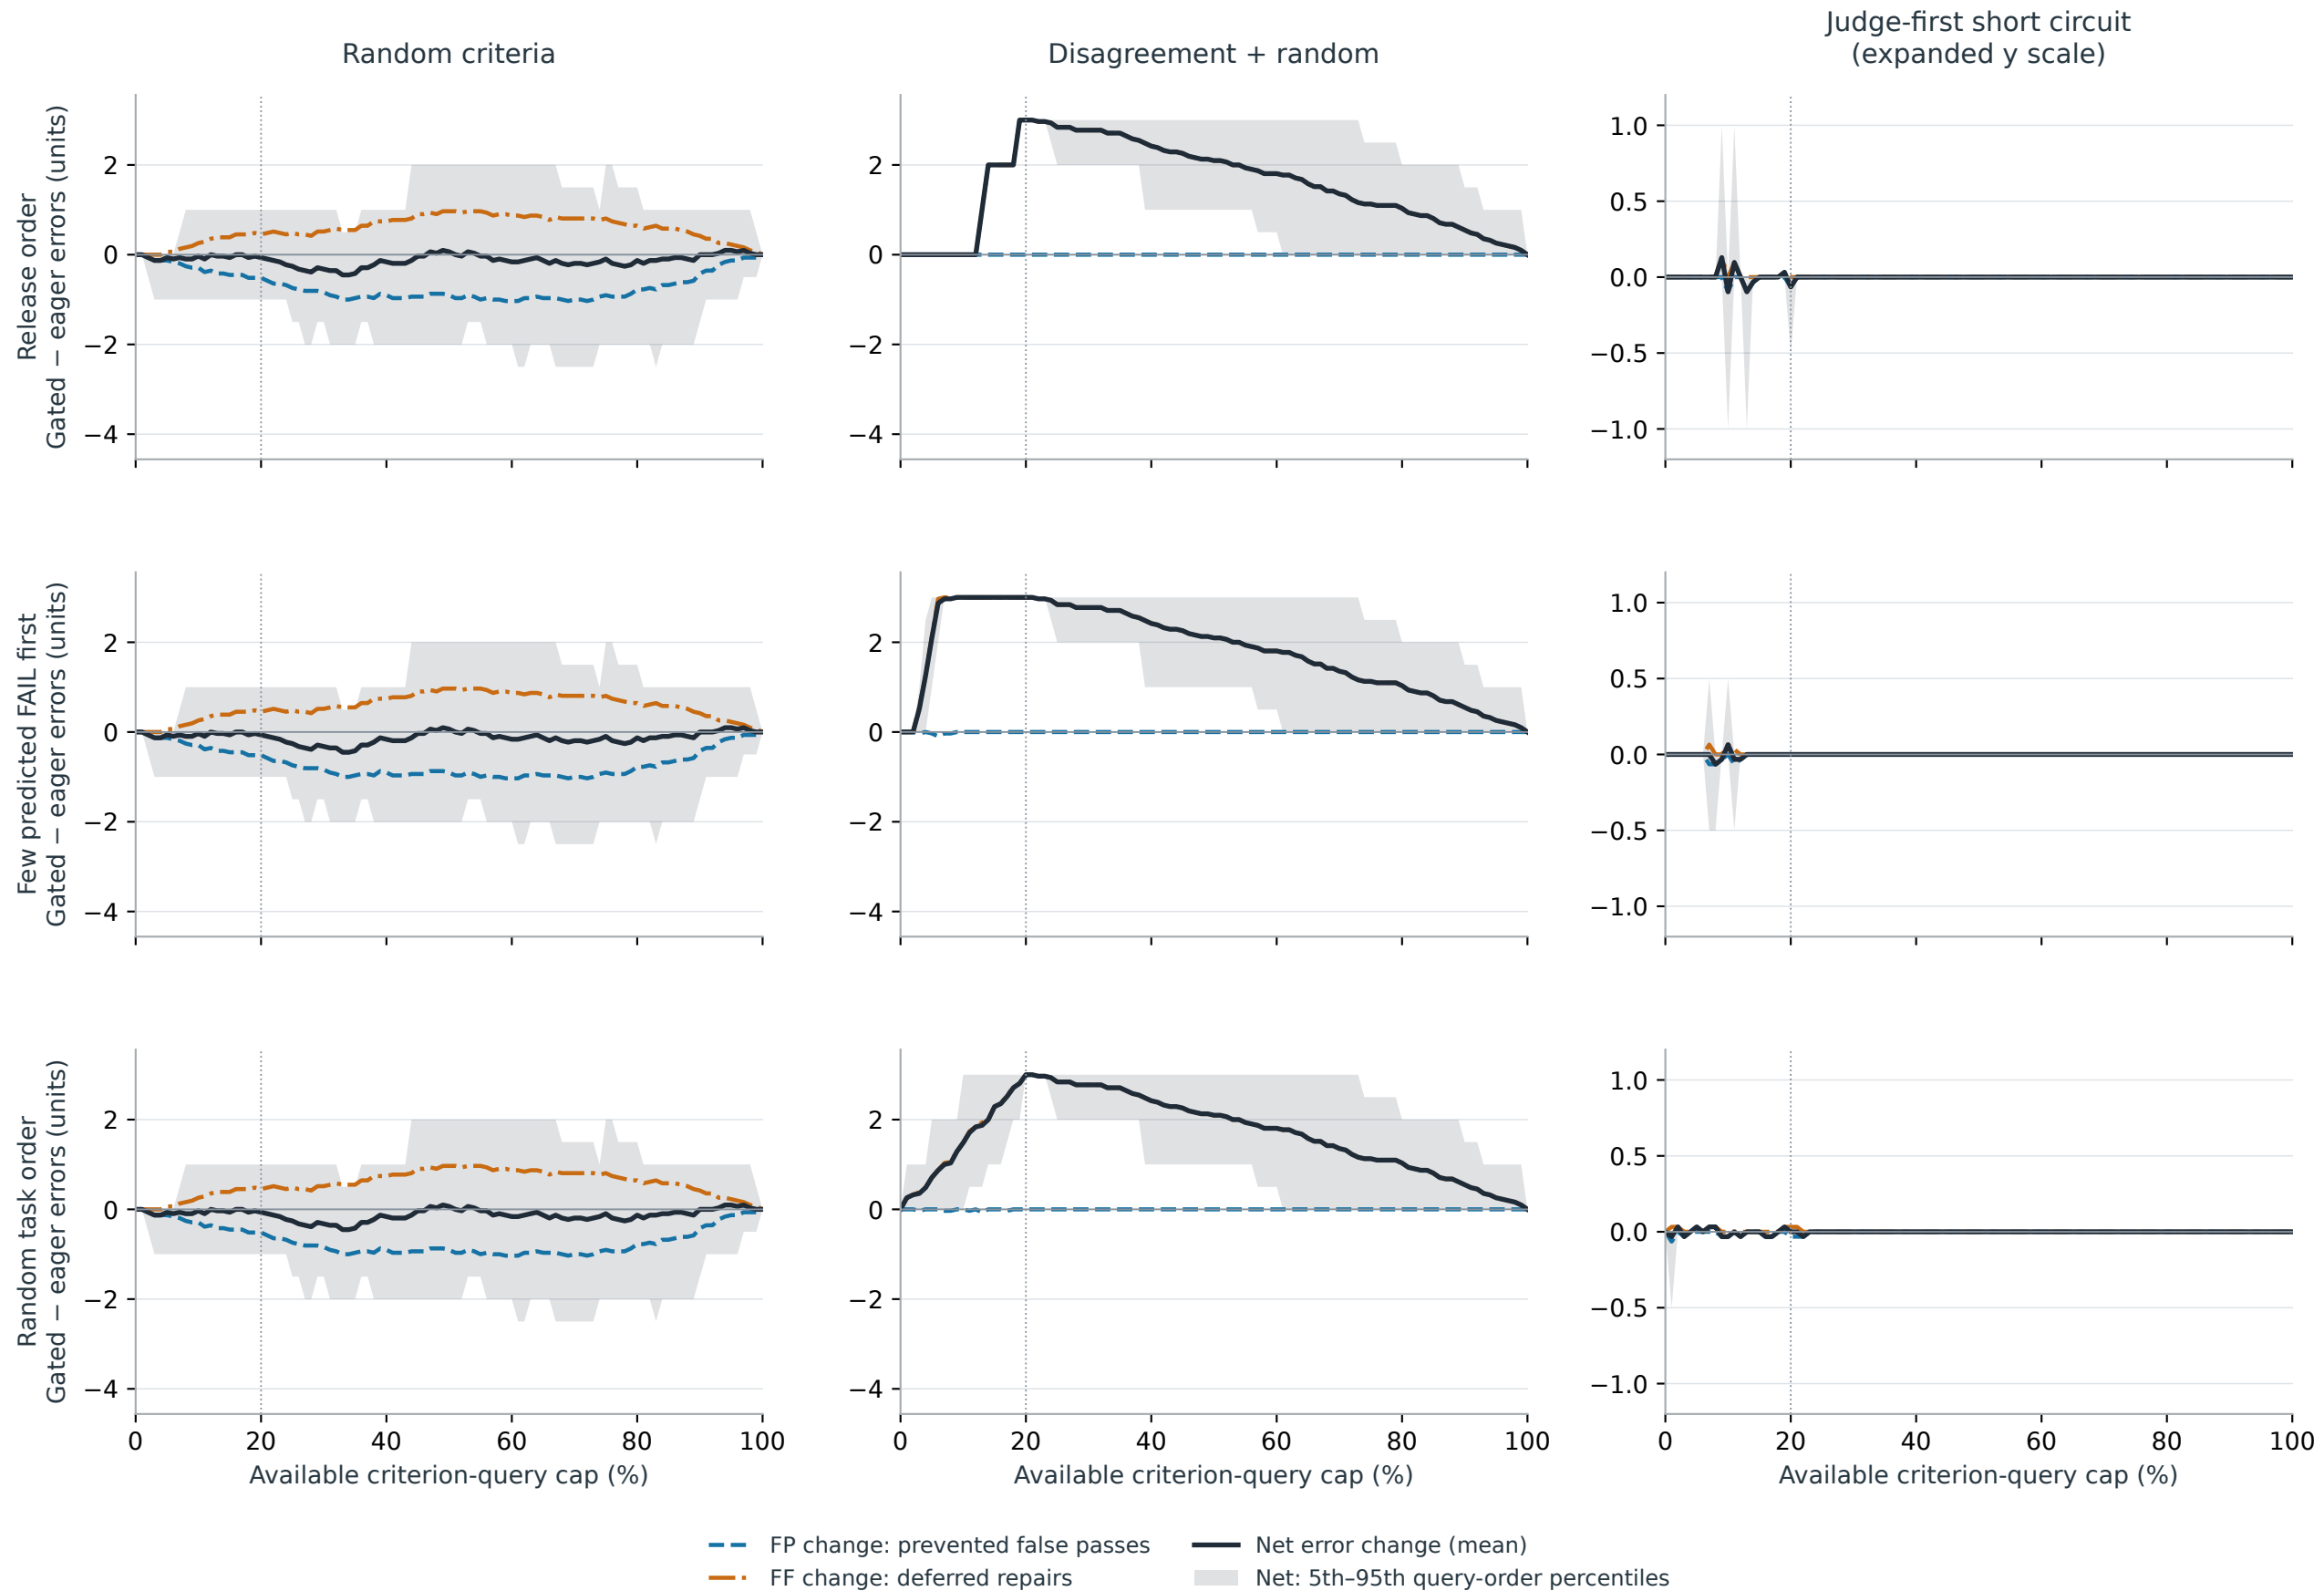
Each point pairs the same queried criteria. Means over seeds 1–31; shaded percentiles are order variation, not confidence intervals.

Negative net: fewer errors with gating. Positive net: more errors with gating. FP + FF = net for means; quantiles do not add.

Budget is an available criterion-query cap; early stopping can leave it unused. Auxiliary judge inference is a separate cost. Fixed released references.

Update effects across all budgets — DR Qwen-plus

N = 284 analysis units; M = 1,615 criteria • shared scale for first two columns within this page; scales vary by page



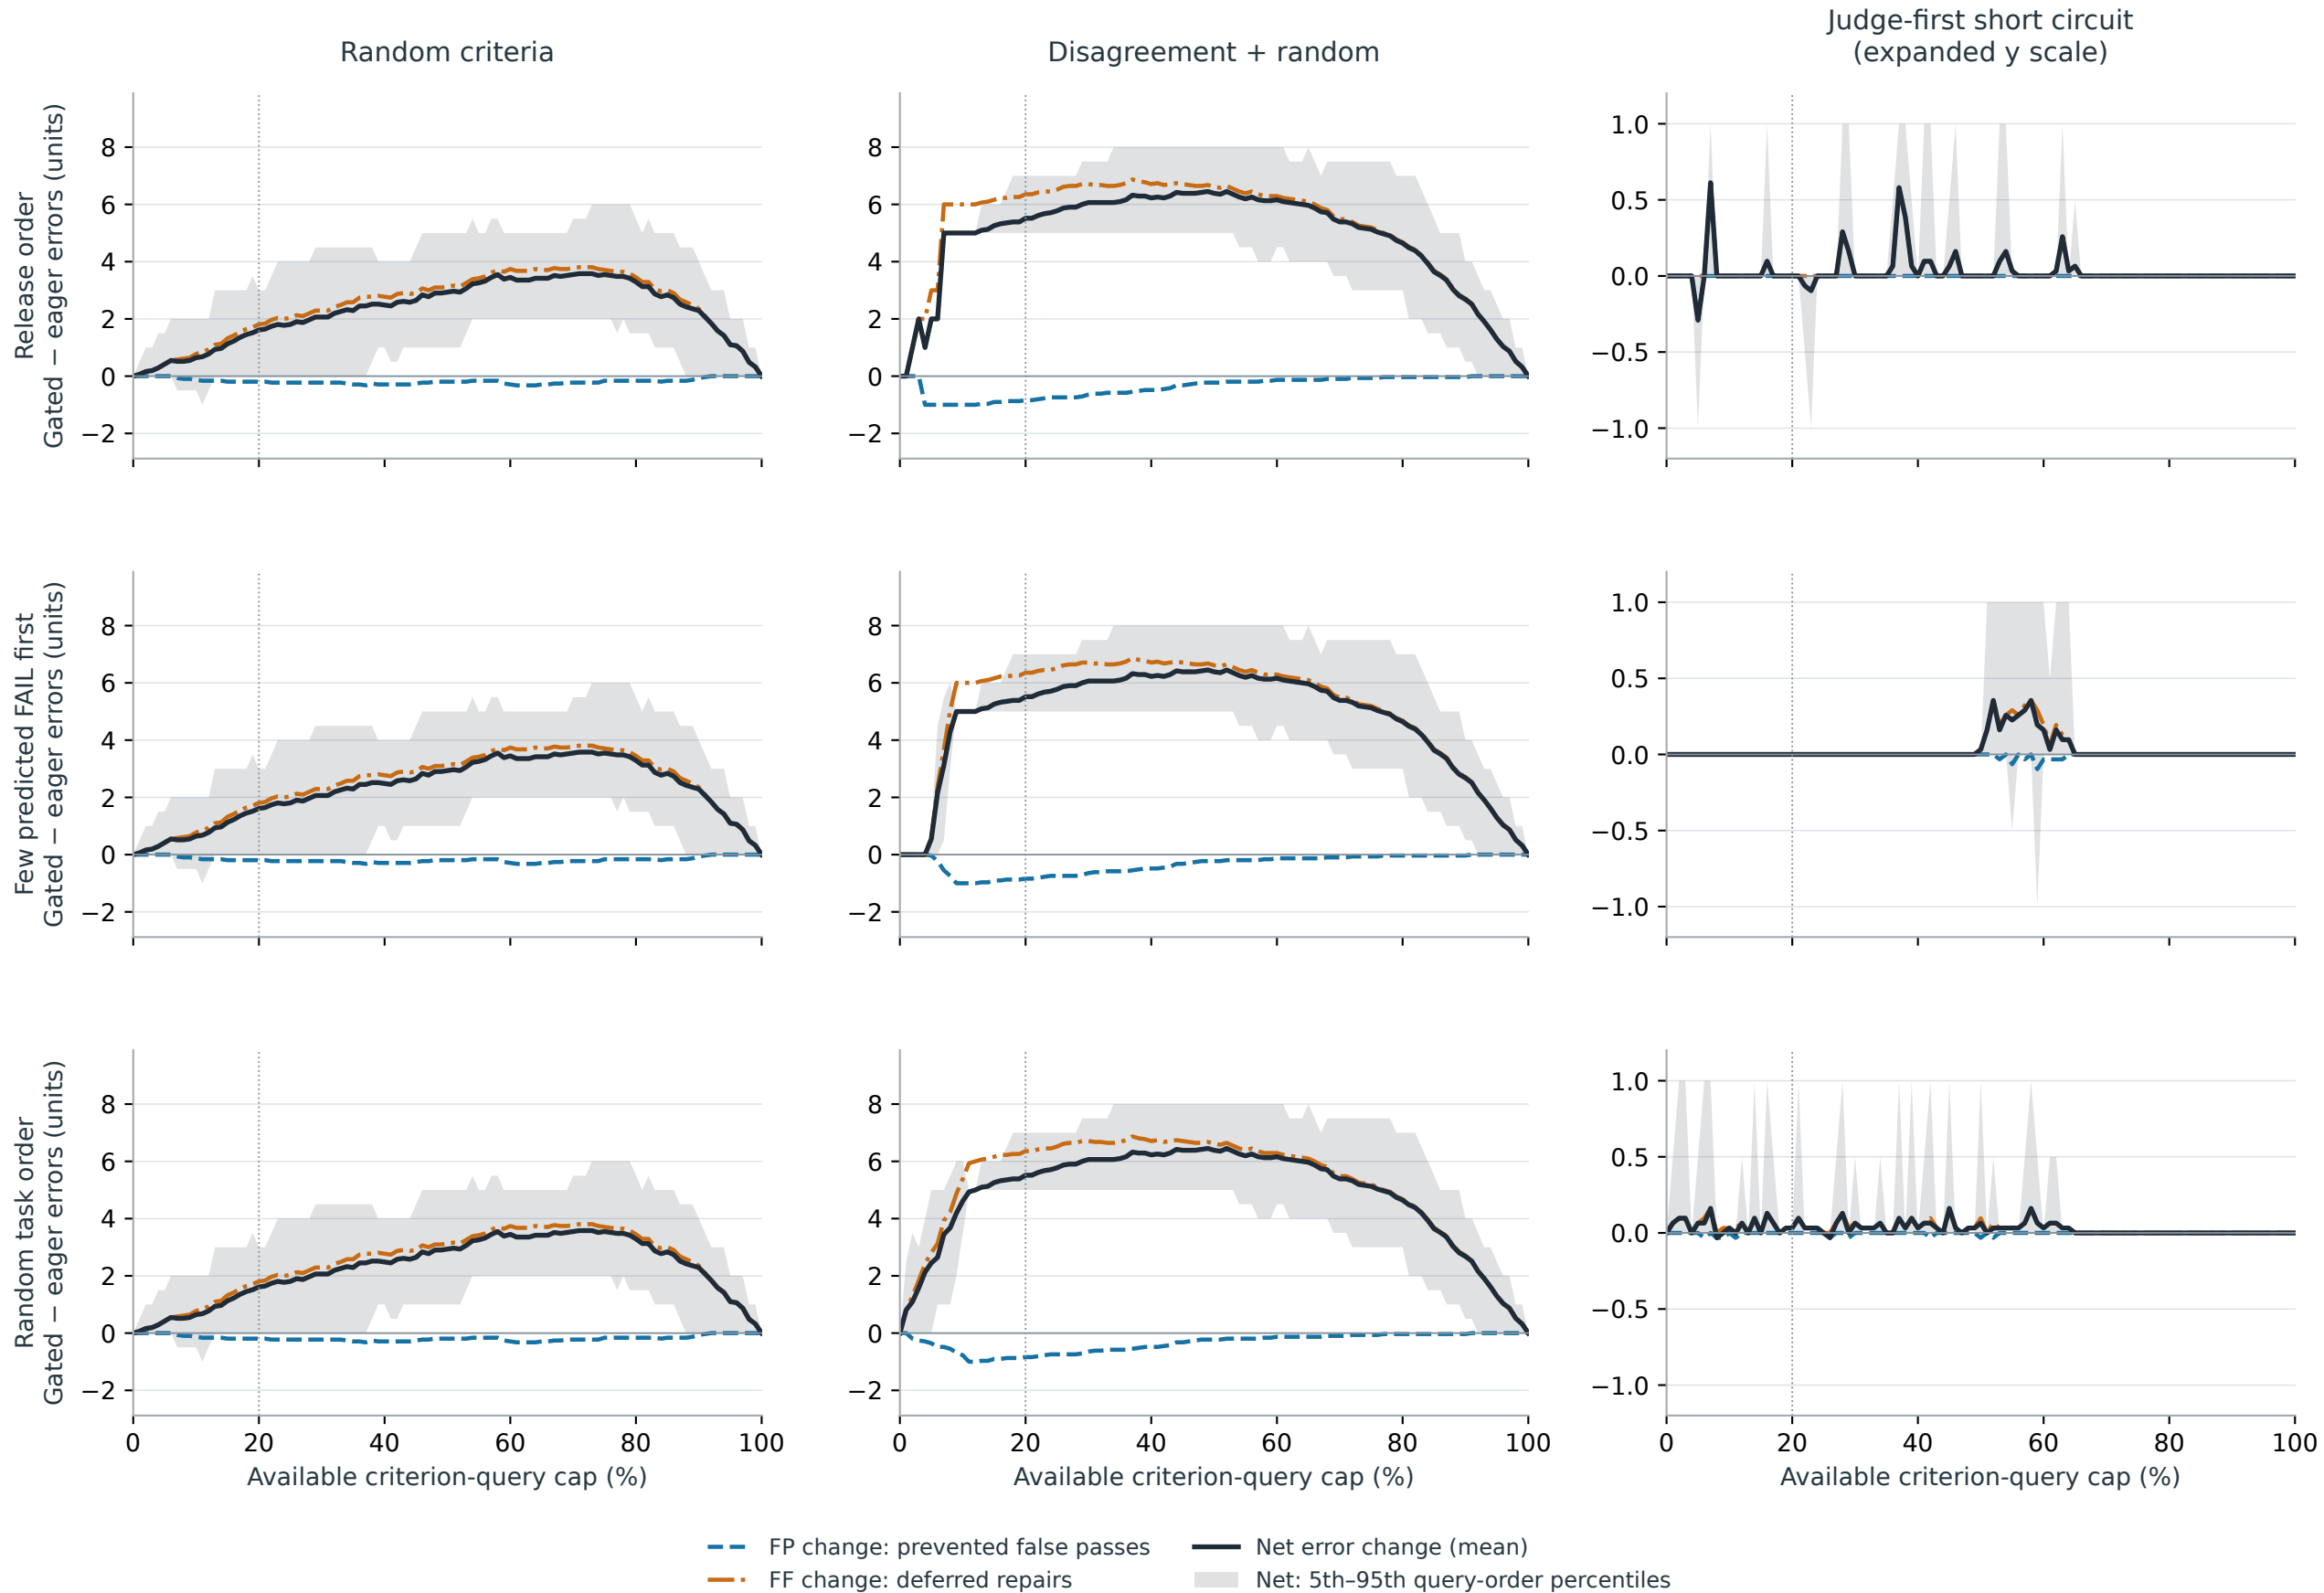
Each point pairs the same queried criteria. Means over seeds 1-31; shaded percentiles are order variation, not confidence intervals.

Negative net: fewer errors with gating. Positive net: more errors with gating. FP + FF = net for means; quantiles do not add.

Budget is an available criterion-query cap; early stopping can leave it unused. Auxiliary judge inference is a separate cost. Fixed released references.

Update effects across all budgets — AC Qwen-plus

N = 210 analysis units; M = 843 criteria • shared scale for first two columns within this page; scales vary by page



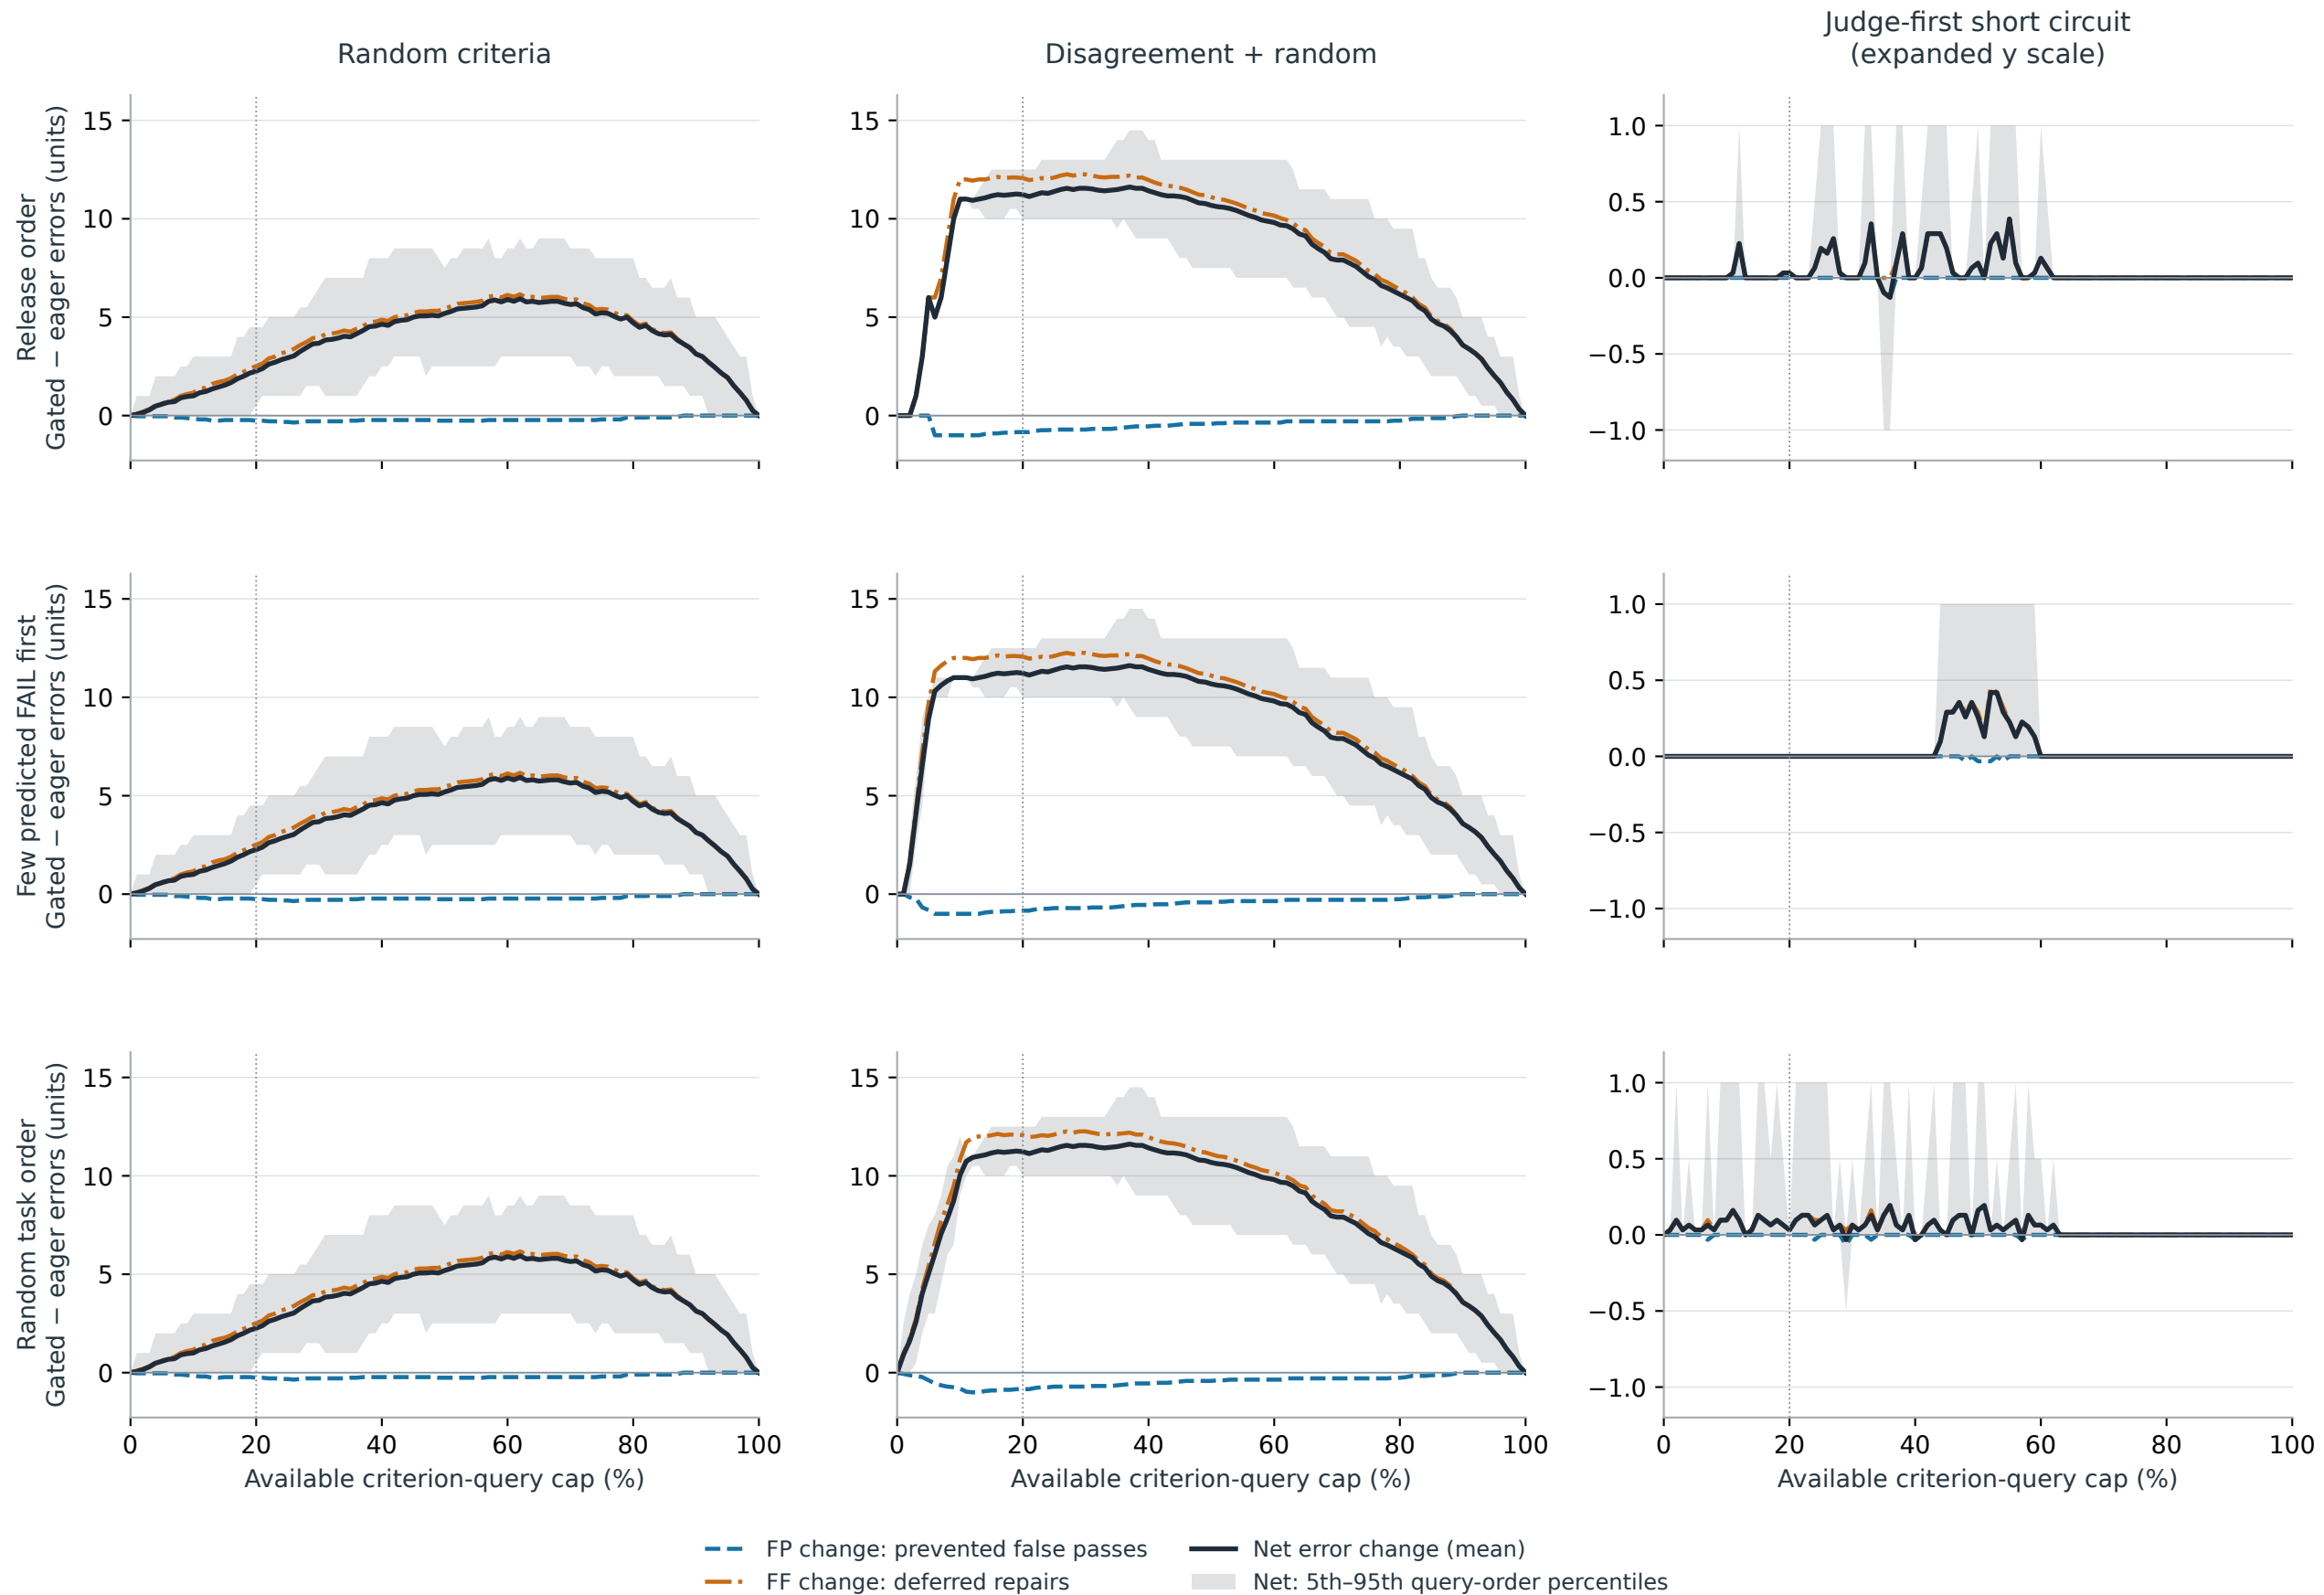
Each point pairs the same queried criteria. Means over seeds 1-31; shaded percentiles are order variation, not confidence intervals.

Negative net: fewer errors with gating. Positive net: more errors with gating. FP + FF = net for means; quantiles do not add.

Budget is an available criterion-query cap; early stopping can leave it unused. Auxiliary judge inference is a separate cost. Fixed released references.

Update effects across all budgets — AC DeepSeek v4-pro

N = 210 analysis units; M = 843 criteria • shared scale for first two columns within this page; scales vary by page



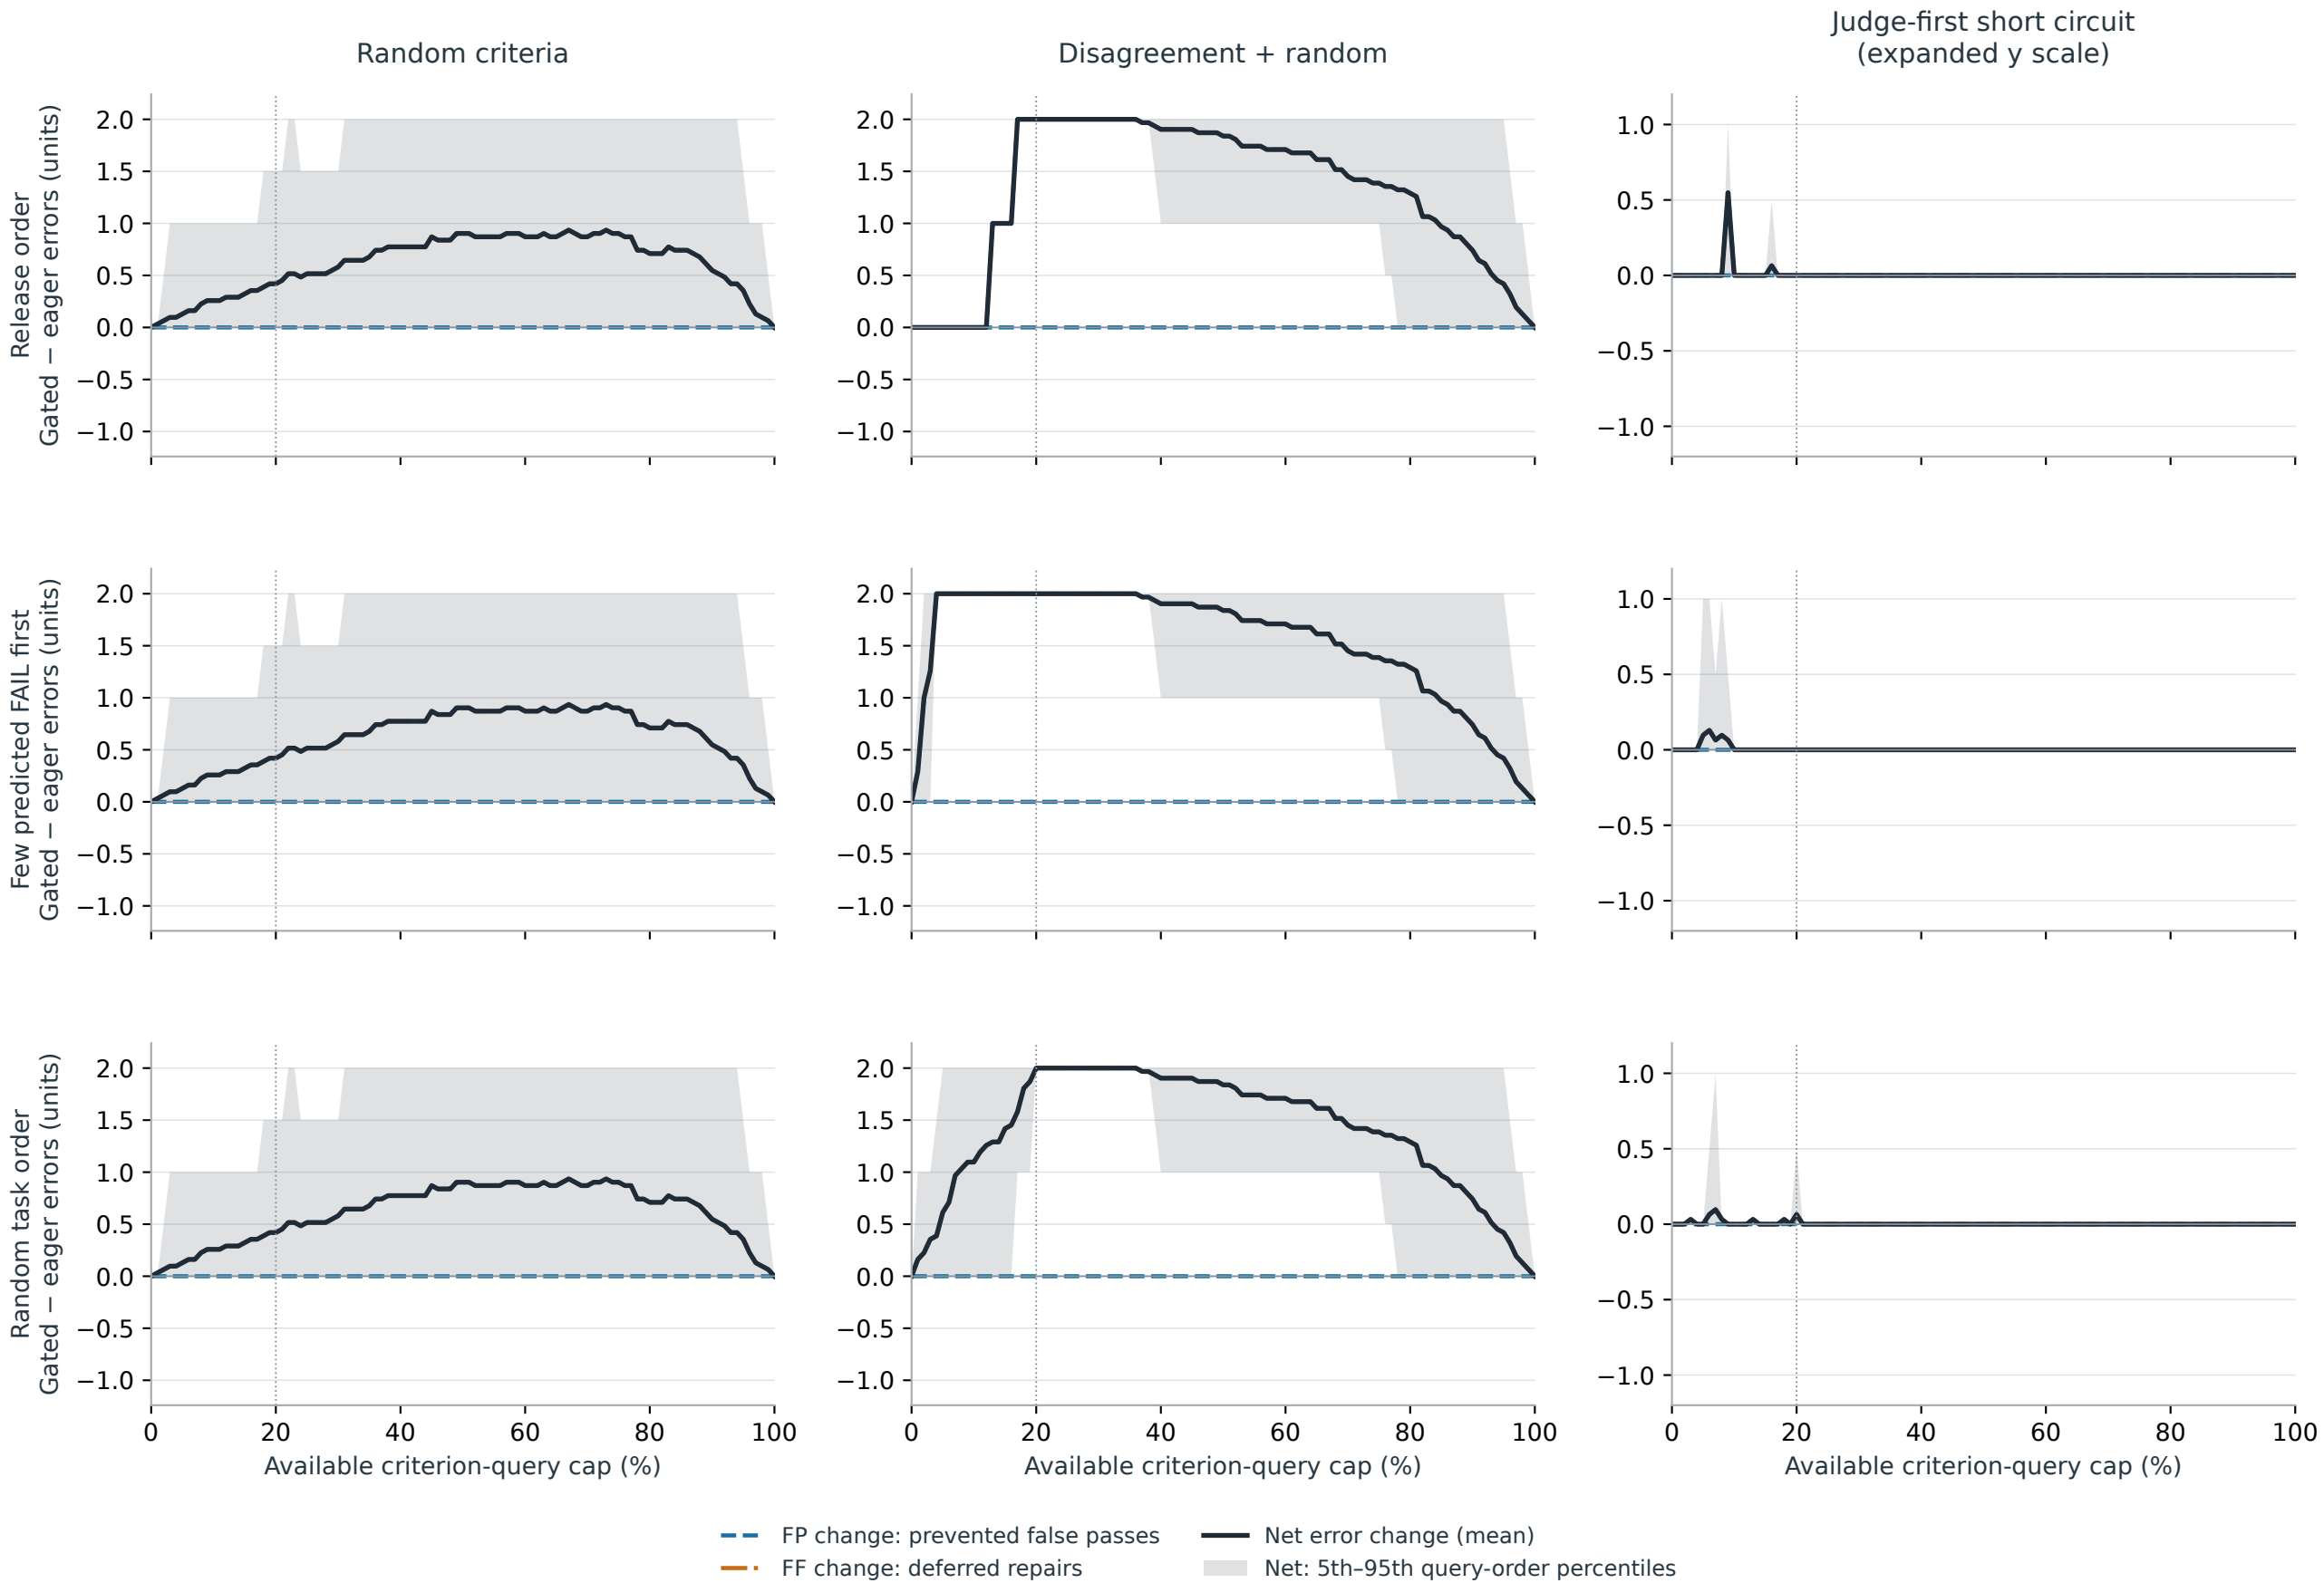
Each point pairs the same queried criteria. Means over seeds 1-31; shaded percentiles are order variation, not confidence intervals.

Negative net: fewer errors with gating. Positive net: more errors with gating. FP + FF = net for means; quantiles do not add.

Budget is an available criterion-query cap; early stopping can leave it unused. Auxiliary judge inference is a separate cost. Fixed released references.

Update effects across all budgets — DR GPT-5.4 low

N = 284 analysis units; M = 1,615 criteria • shared scale for first two columns within this page; scales vary by page



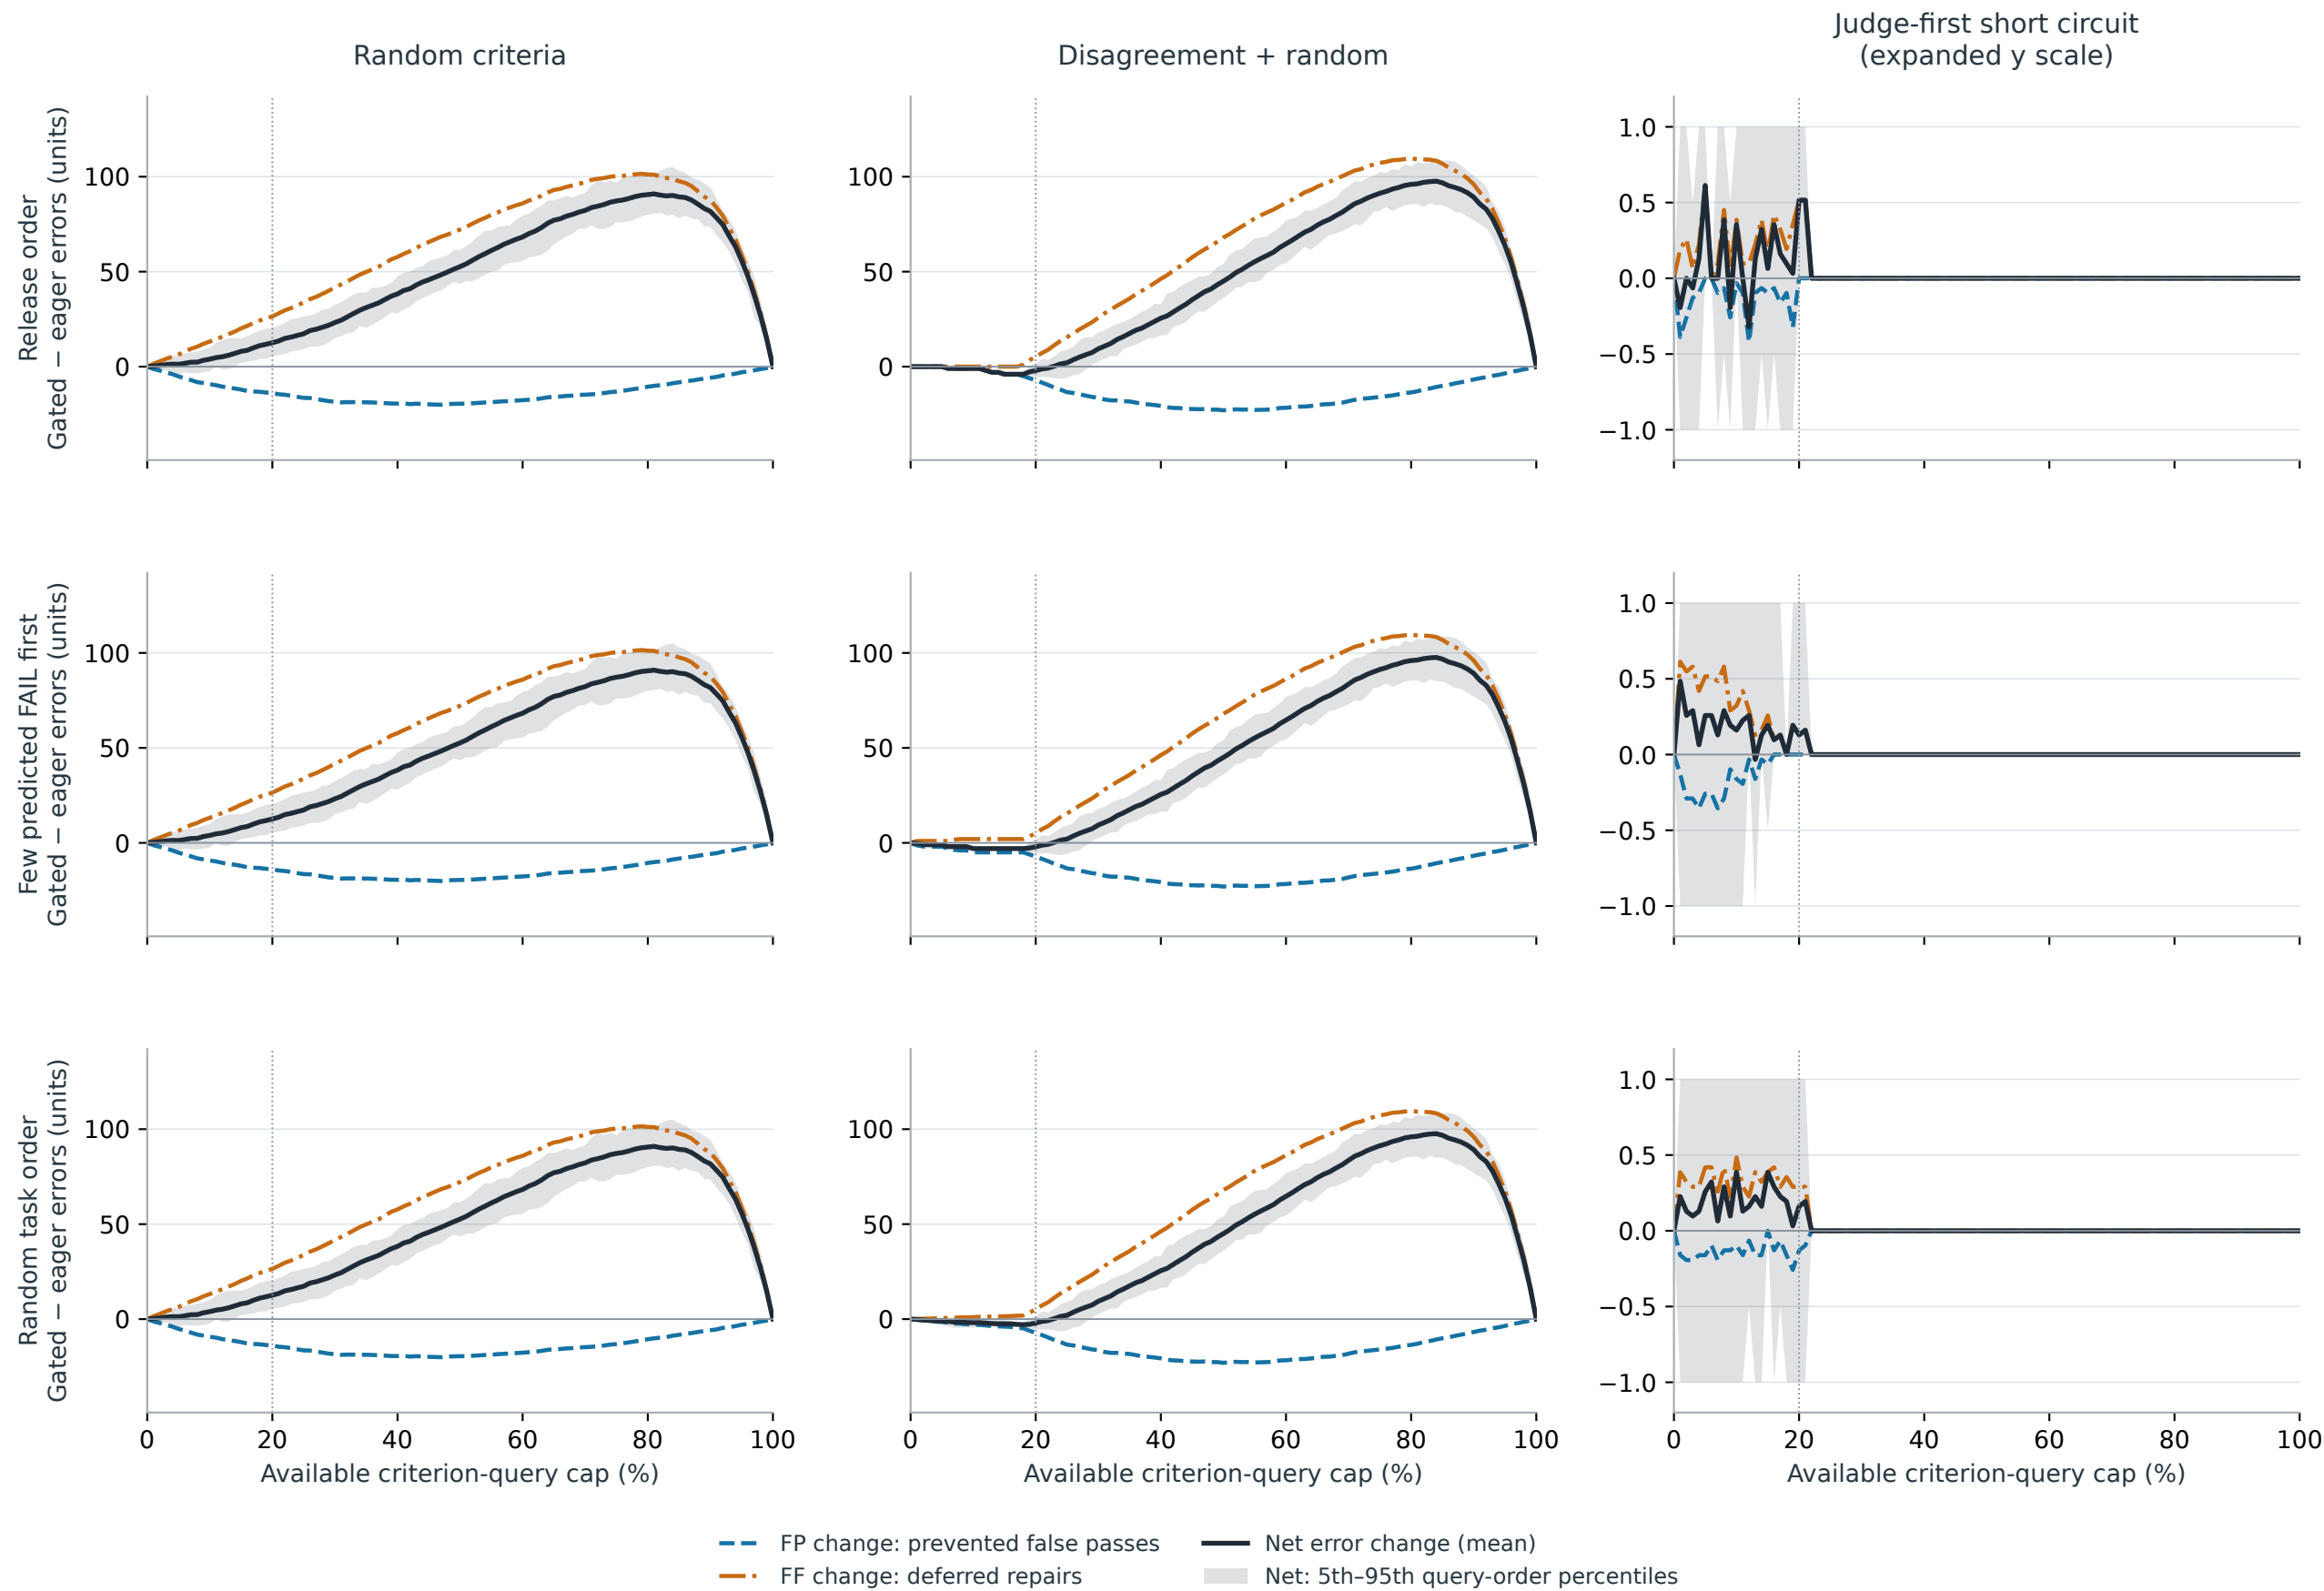
Each point pairs the same queried criteria. Means over seeds 1-31; shaded percentiles are order variation, not confidence intervals.

Negative net: fewer errors with gating. Positive net: more errors with gating. FP + FF = net for means; quantiles do not add.

Budget is an available criterion-query cap; early stopping can leave it unused. Auxiliary judge inference is a separate cost. Fixed released references.

Update effects across all budgets — JB GPT-5.4

N = 1,539 analysis units; M = 23,487 criteria • shared scale for first two columns within this page; scales vary by page



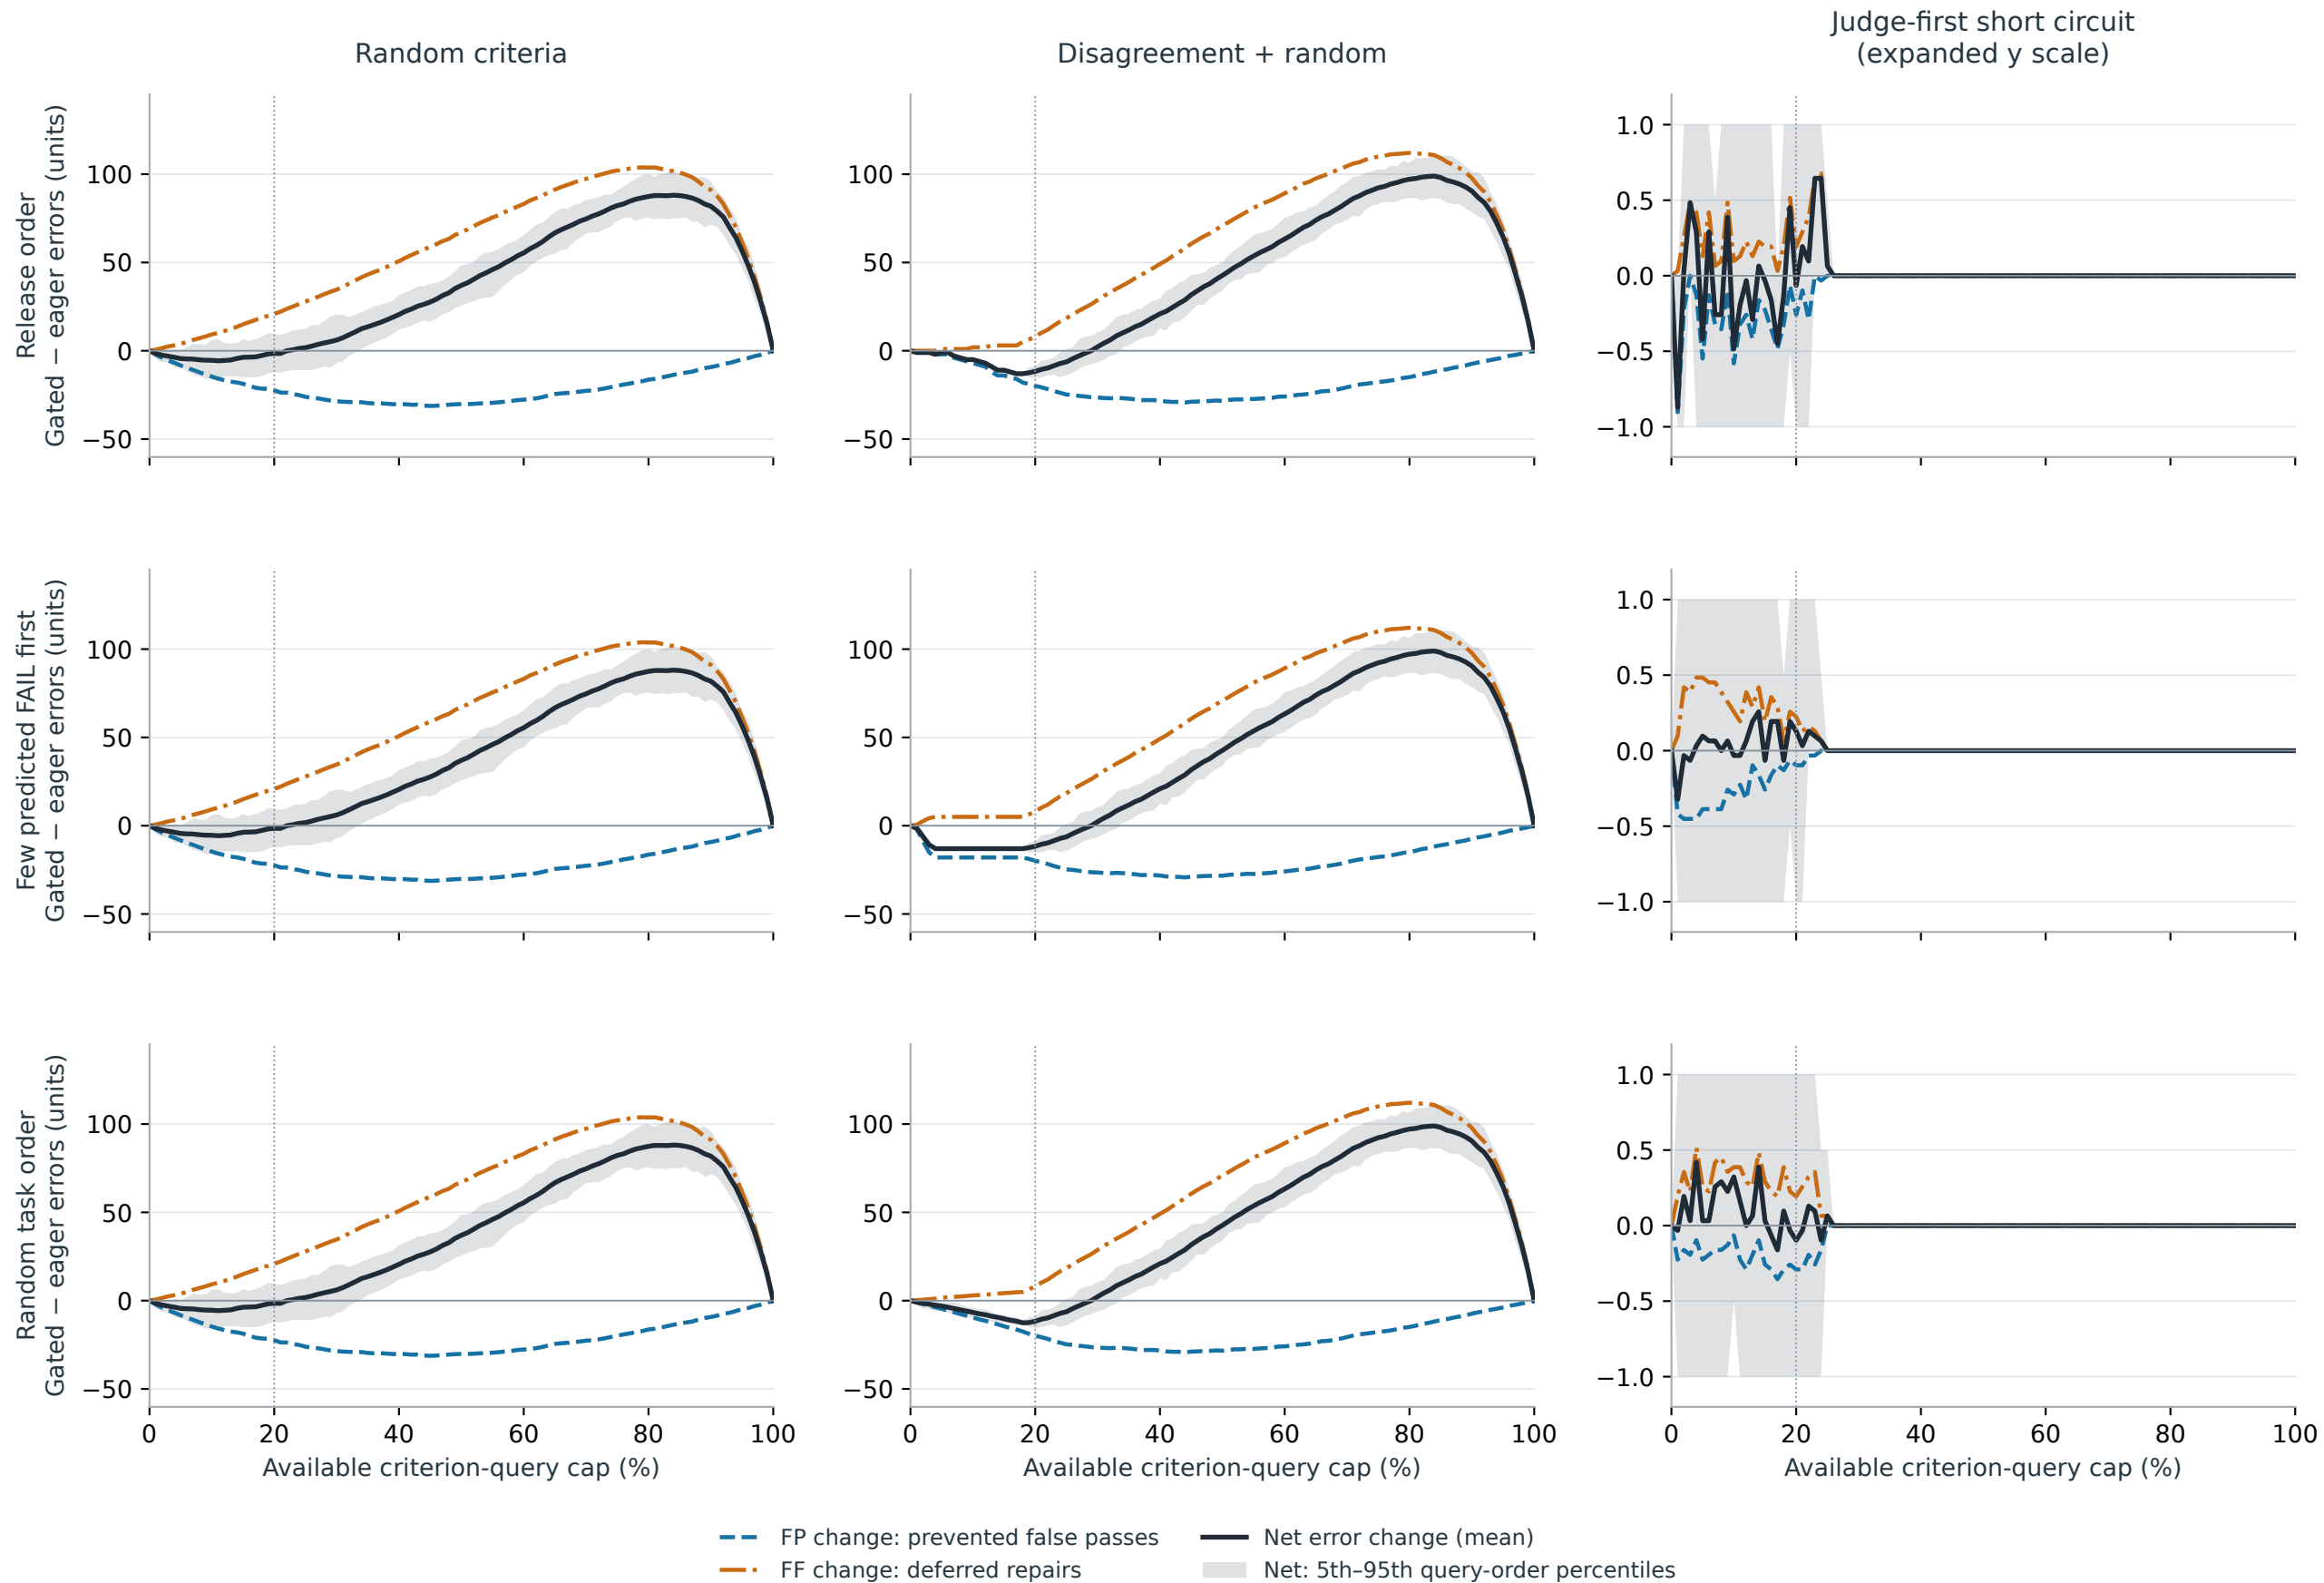
Each point pairs the same queried criteria. Means over seeds 1-31; shaded percentiles are order variation, not confidence intervals.

Negative net: fewer errors with gating. Positive net: more errors with gating. FP + FF = net for means; quantiles do not add.

Budget is an available criterion-query cap; early stopping can leave it unused. Auxiliary judge inference is a separate cost. Fixed released references.

Update effects across all budgets — JB GPT-5.4-mini

N = 1,539 analysis units; M = 23,487 criteria • shared scale for first two columns within this page; scales vary by page



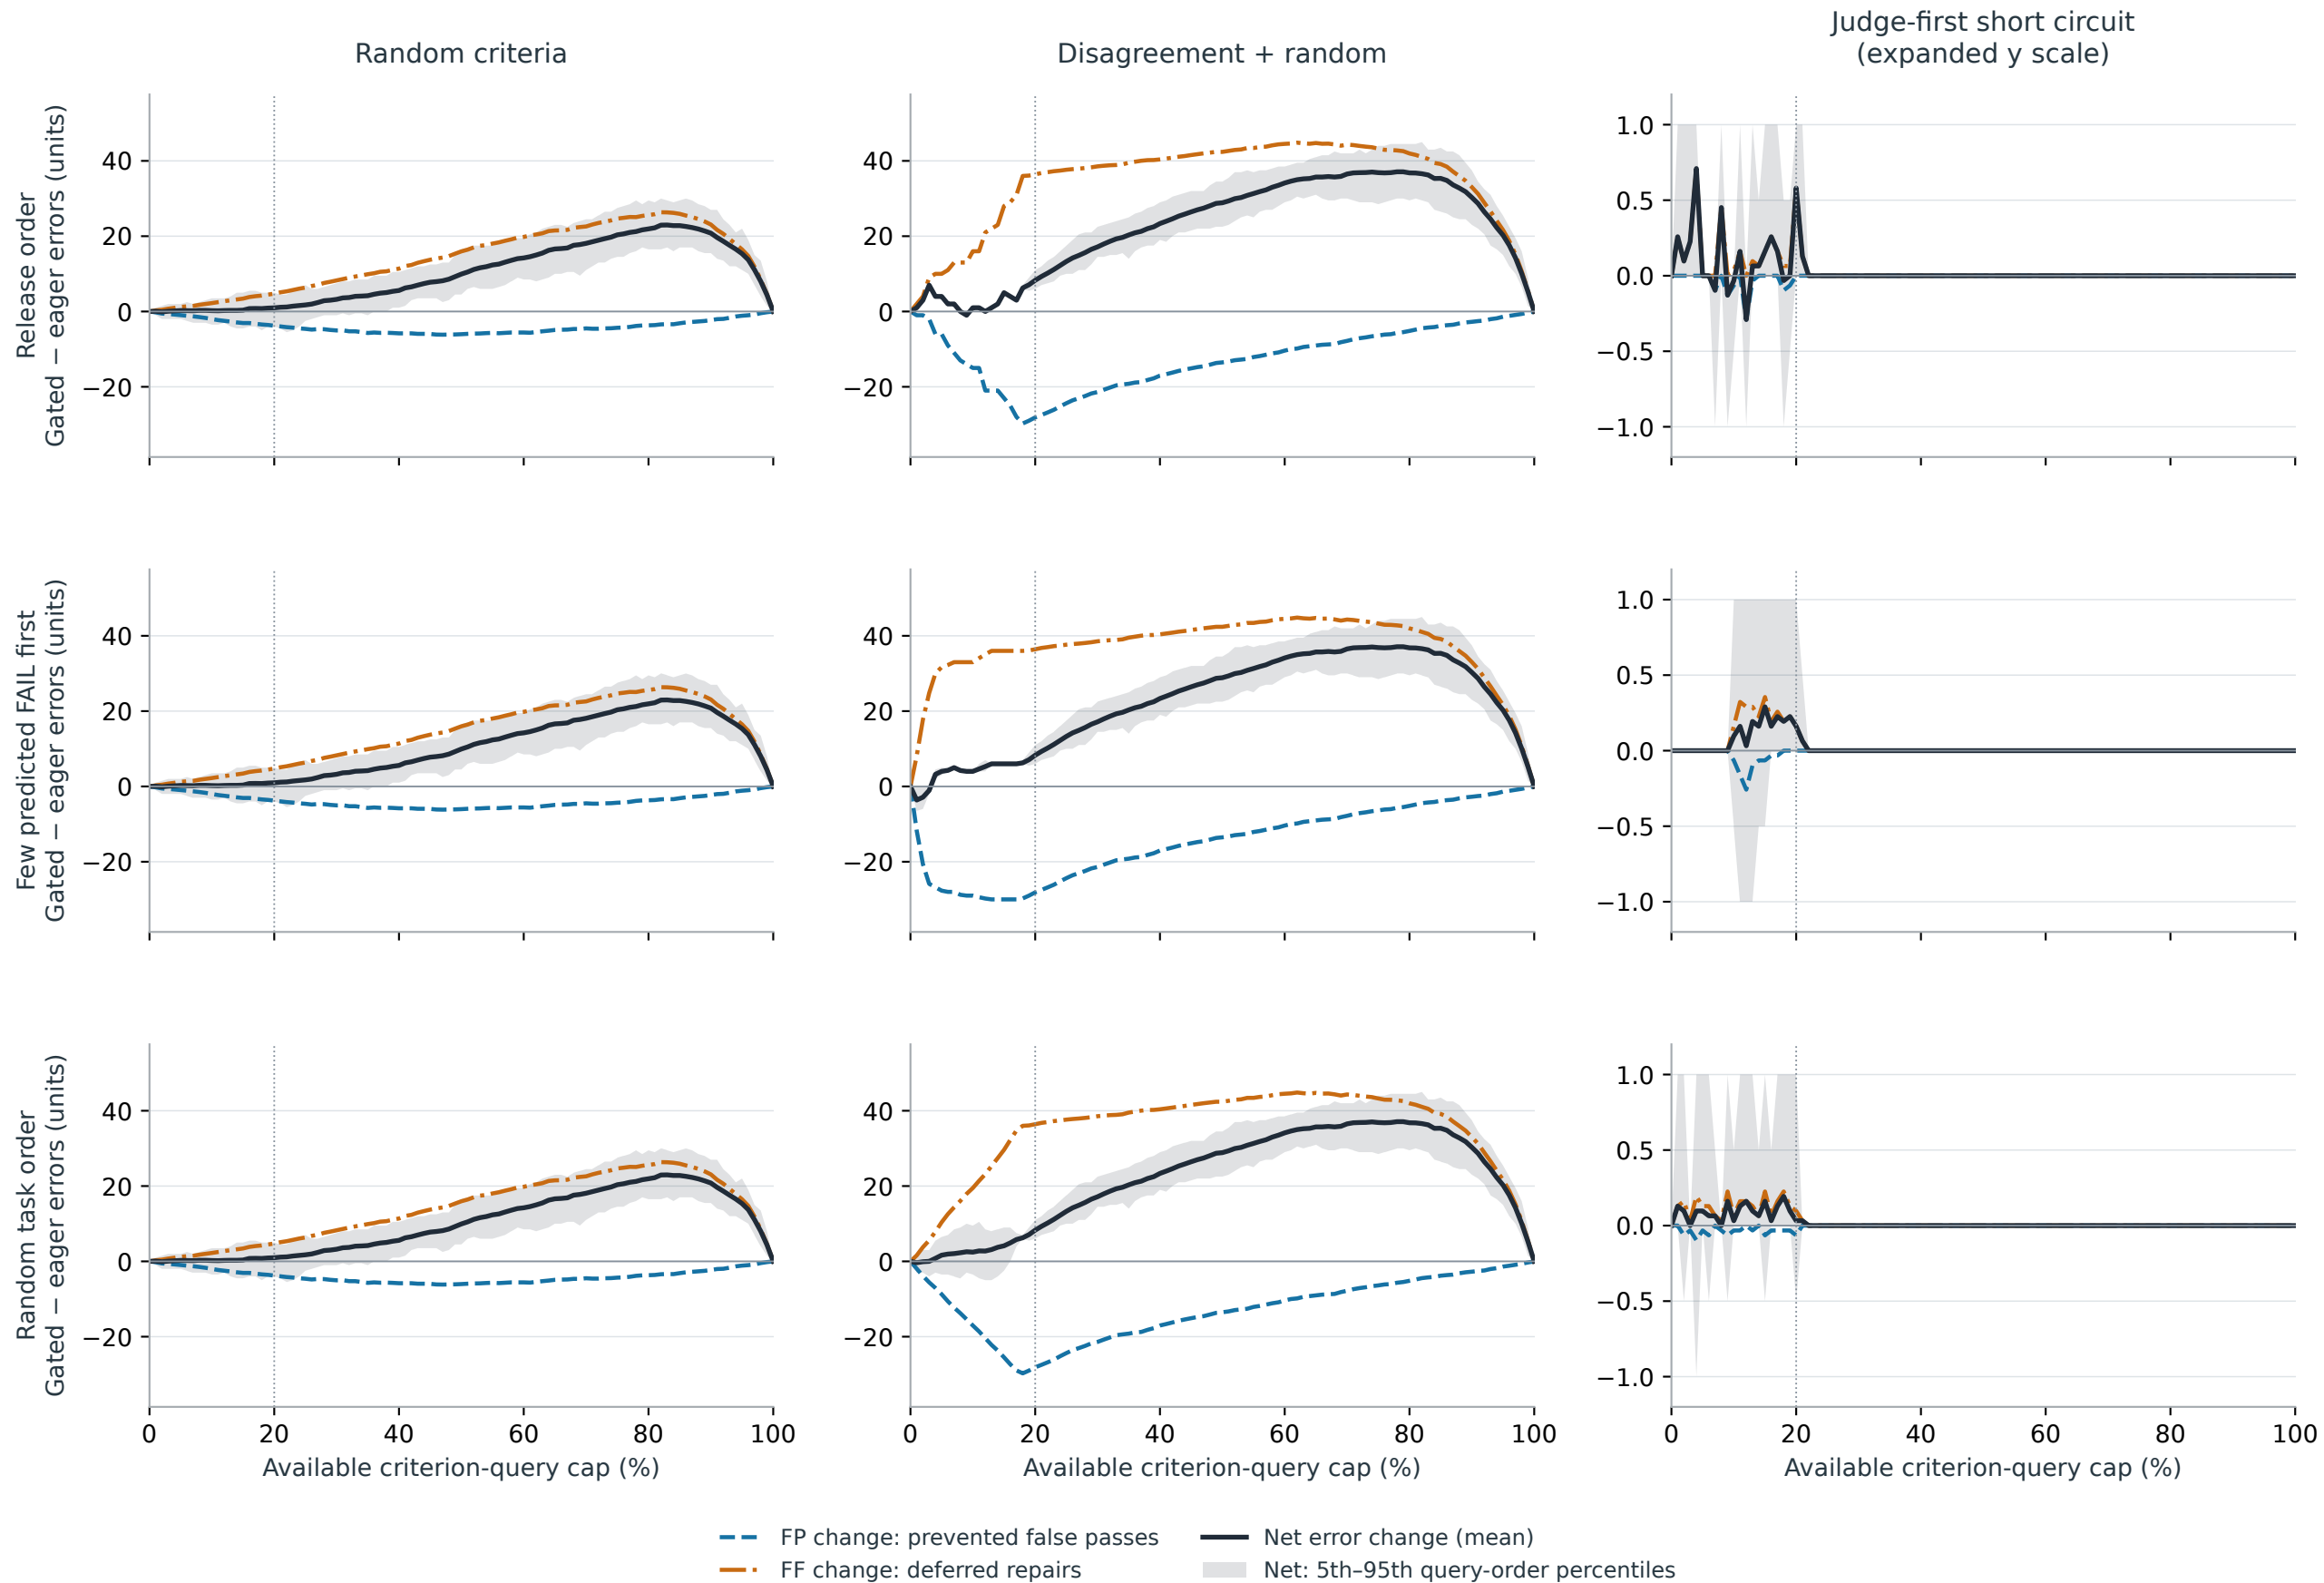
Each point pairs the same queried criteria. Means over seeds 1–31; shaded percentiles are order variation, not confidence intervals.

Negative net: fewer errors with gating. Positive net: more errors with gating. FP + FF = net for means; quantiles do not add.

Budget is an available criterion-query cap; early stopping can leave it unused. Auxiliary judge inference is a separate cost. Fixed released references.

Update effects across all budgets — JB GPT-5.4 | binary_only

N = 1,539 analysis units; M = 20,229 criteria • shared scale for first two columns within this page; scales vary by page



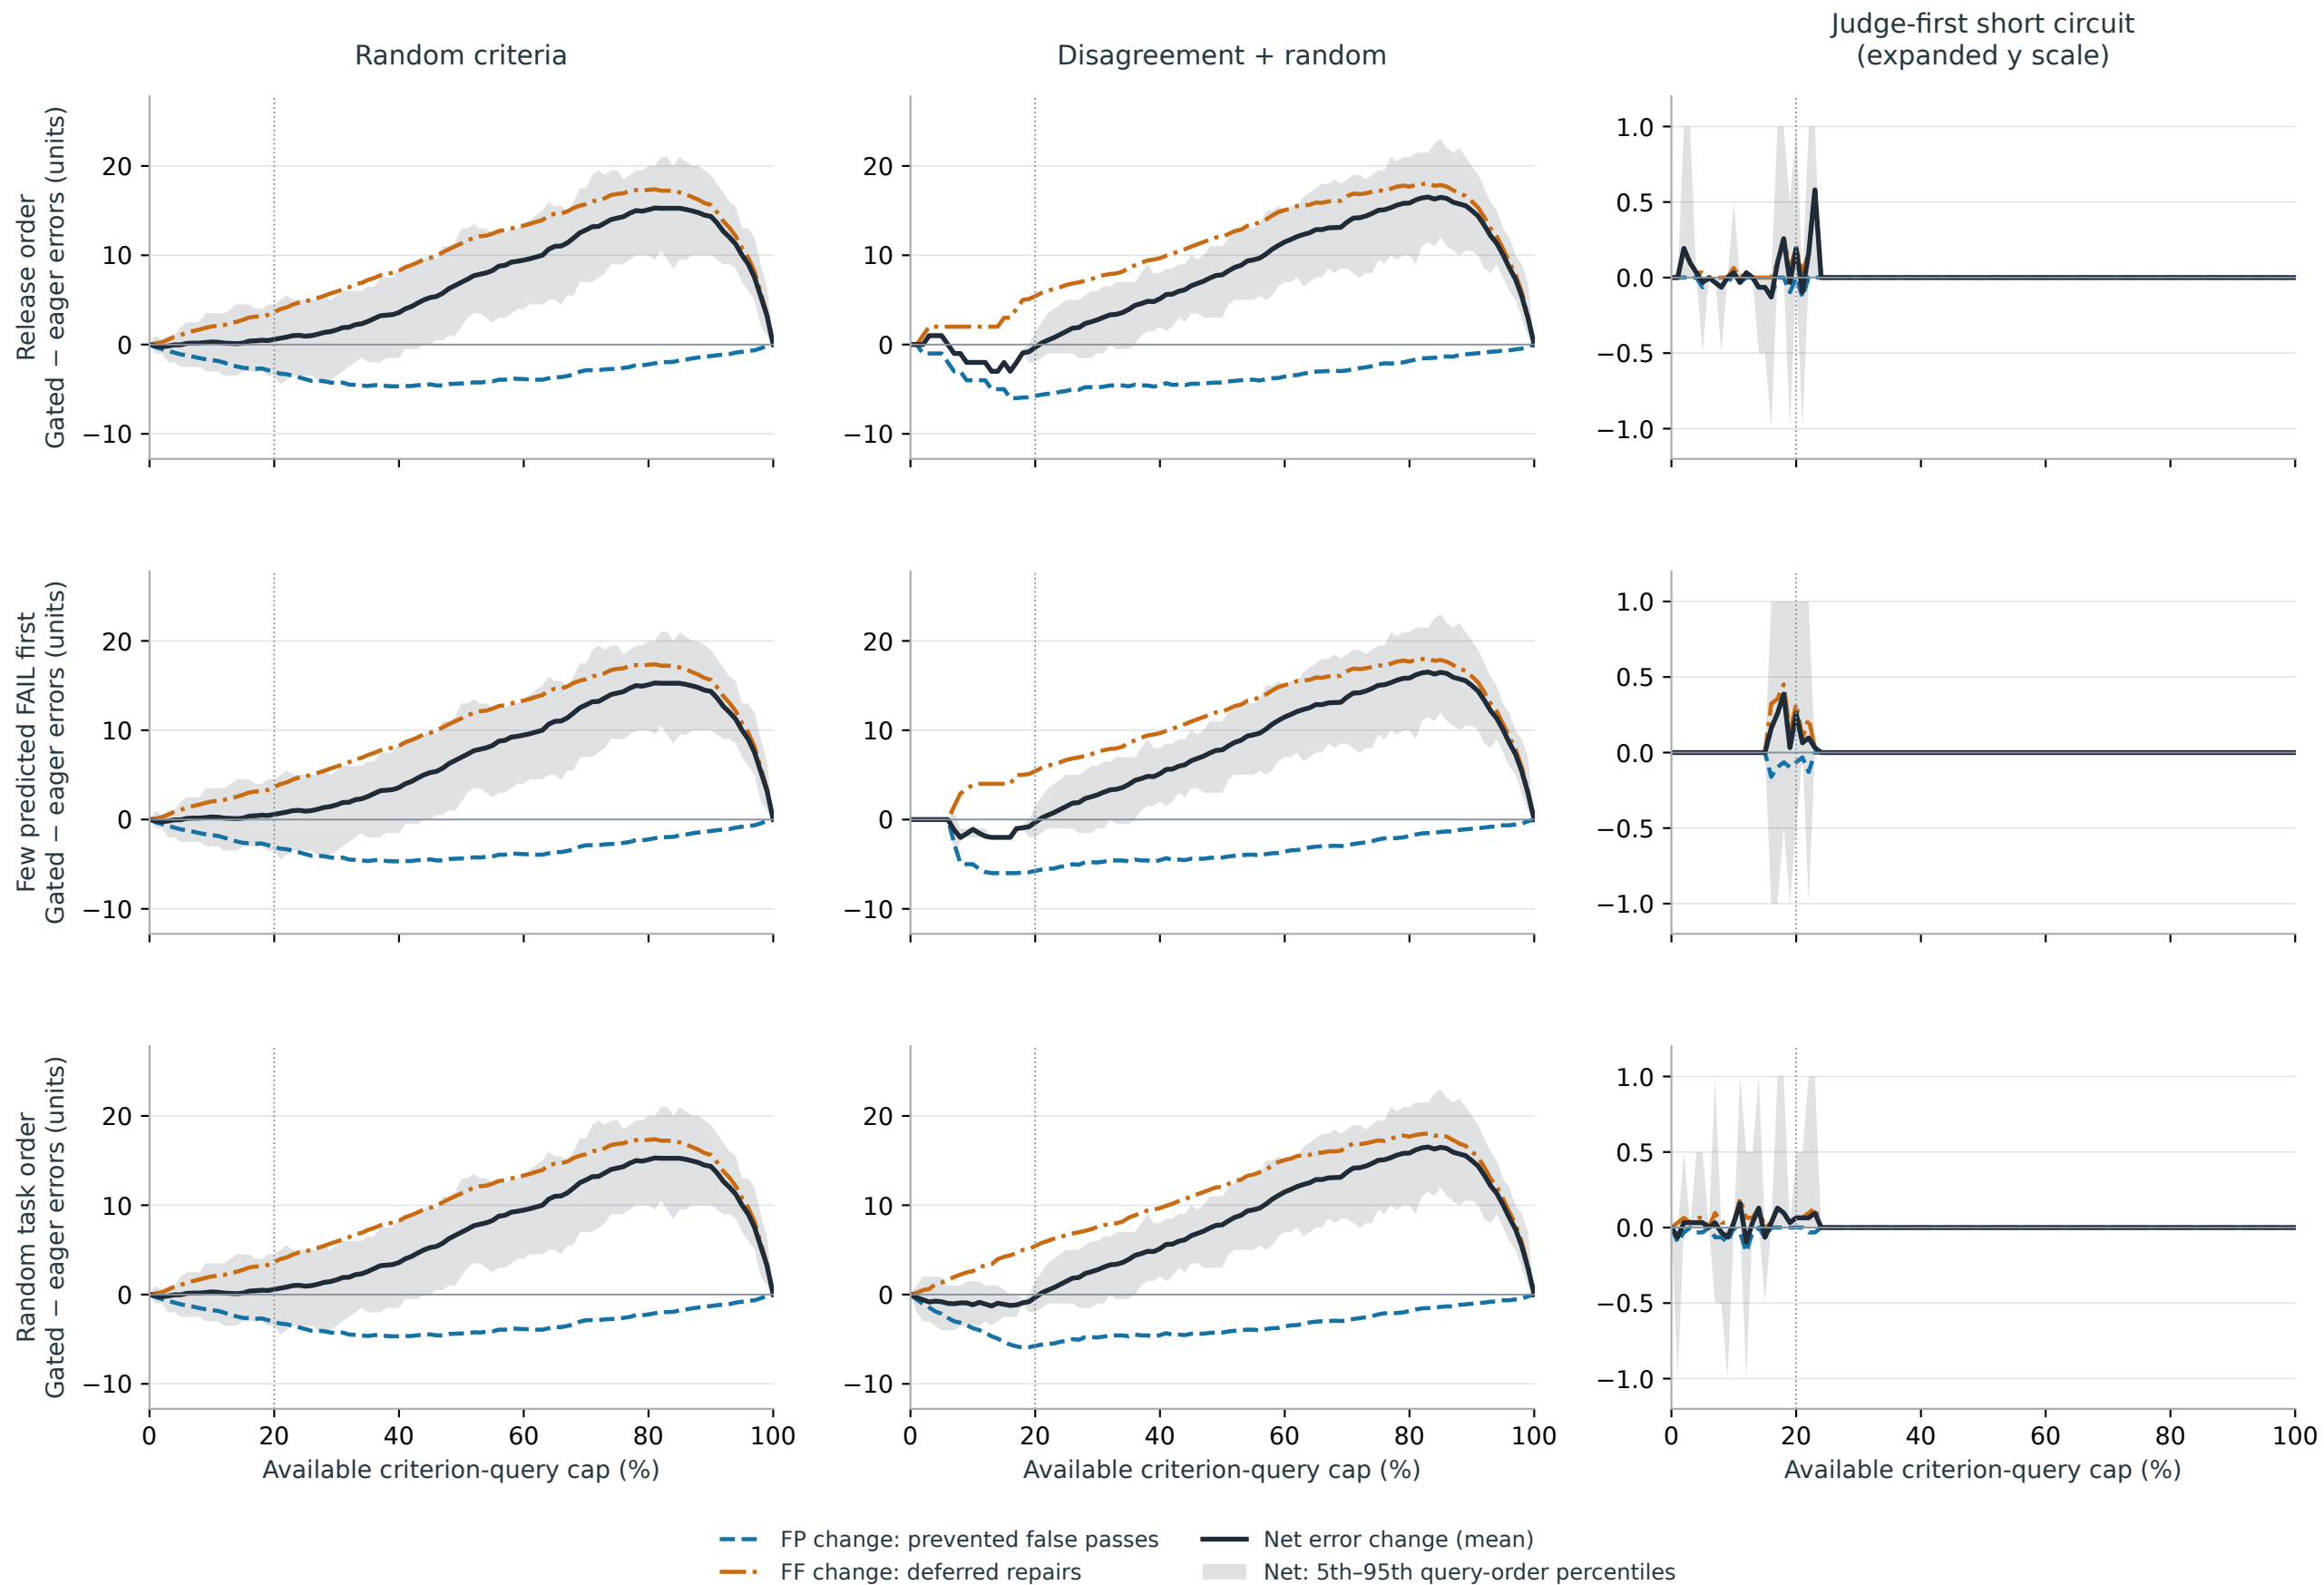
Each point pairs the same queried criteria. Means over seeds 1-31; shaded percentiles are order variation, not confidence intervals.

Negative net: fewer errors with gating. Positive net: more errors with gating. FP + FF = net for means; quantiles do not add.

Budget is an available criterion-query cap; early stopping can leave it unused. Auxiliary judge inference is a separate cost. Fixed released references.

Update effects across all budgets — JB GPT-5.4-mini | binary_only

N = 1,539 analysis units; M = 20,229 criteria • shared scale for first two columns within this page; scales vary by page



Each point pairs the same queried criteria. Means over seeds 1-31; shaded percentiles are order variation, not confidence intervals.

Negative net: fewer errors with gating. Positive net: more errors with gating. FP + FF = net for means; quantiles do not add.

Budget is an available criterion-query cap; early stopping can leave it unused. Auxiliary judge inference is a separate cost. Fixed released references.