Why verdict gating can help or hurt at the same query budget

Full-reference expectation | disagreement first, random continuation | release order

(a) JudgmentBench: GPT-5.4, full

N = 1,539 annotations; M = 23,487 criteria

(b) JudgmentBench · GPT-5.4 · binary only

N = 1,539 annotations; M = 20,229 criteria

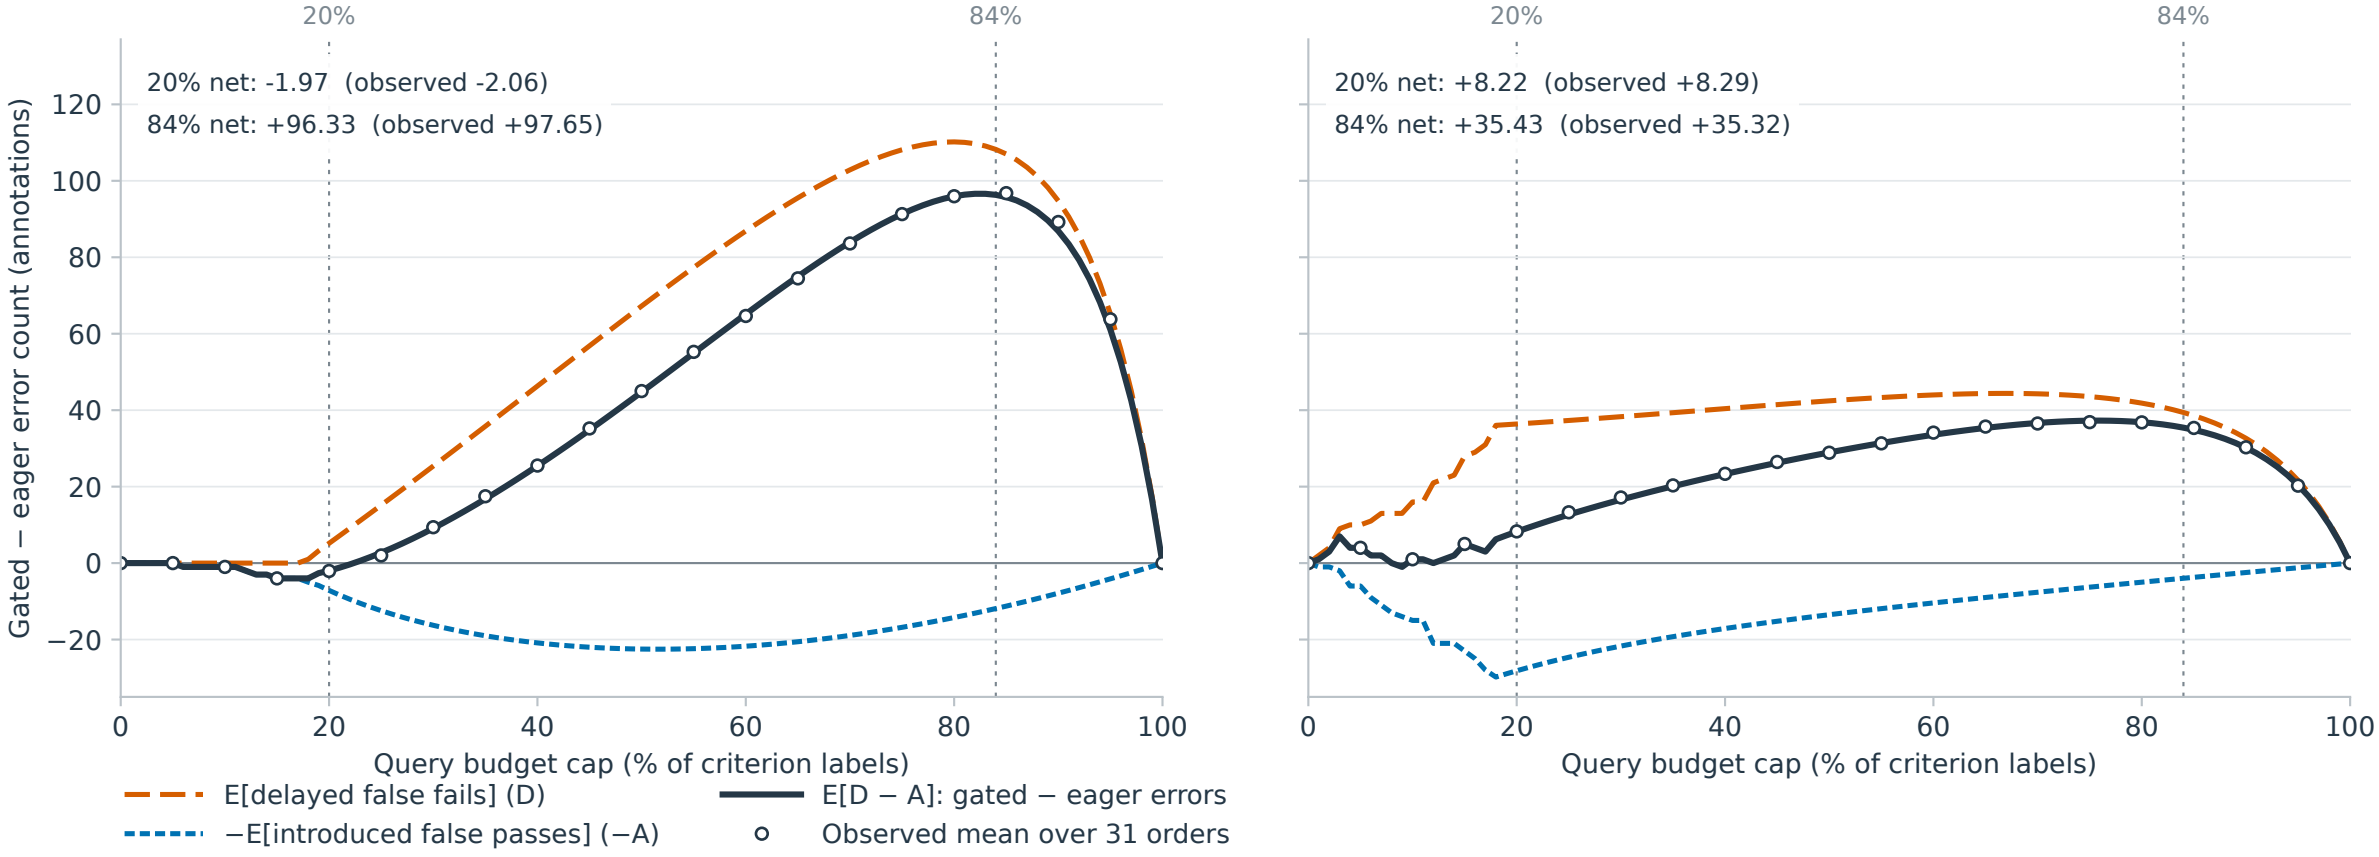

Positive net = more errors under gating. Expectations condition on all released reference labels; they are retrospective, not deployable forecasts.
Points average seeds 1–31 (seed 0 excluded), every 5 percentage points. These order means do not quantify new-task uncertainty.